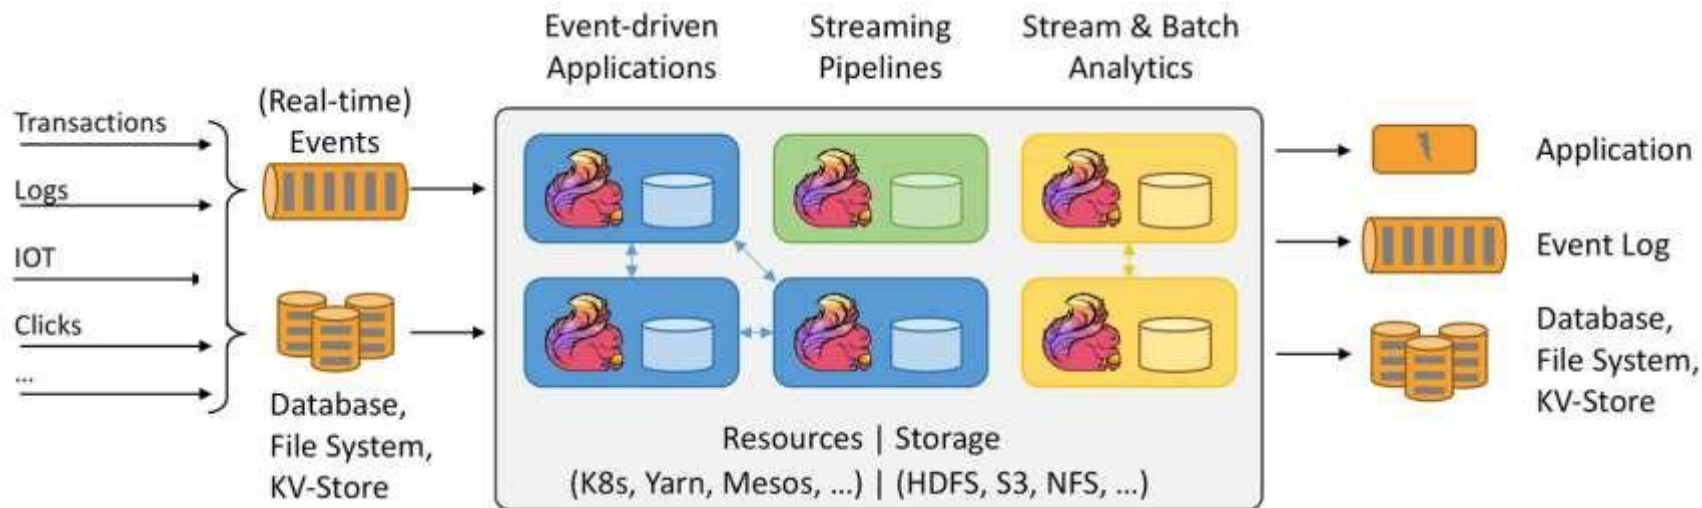


# Apache Flink<sup>®</sup> SQL

Based on slides by Fabian Hueske and others

# What is Apache Flink?

Stateful computations over streams  
real-time and historic  
fast, scalable, fault tolerant, in-memory  
event time, large state, exactly-once



# Flink's Powerful Abstractions

Layered abstractions to  
navigate simple to complex use cases

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp)
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```

High-level  
Analytics API

**SQL / Table API (dynamic tables)**



Stream- & Batch  
Data Processing

**DataStream API (streams, windows)**



```
val stats = stream
  .keyBy("sensor")
  .timeWindow(Time.seconds(5))
  .sum((a, b) -> a.add(b))
```

Stateful Event-  
Driven Applications

**Process Function (events, state, time)**



```
def processElement(event: MyEvent, ctx: Context, out: Collector[Result]) = {
  // work with event and state
  (event, state.value) match { ... }

  out.collect(...) // emit events
  state.update(...) // modify state

  // schedule a timer callback
  ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
}
```

# Flink's Relational APIs

## ANSI SQL

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

## LINQ-style Table API

```
tableEnvironment  
    .scan("clicks")  
    .groupBy('user')  
    .select('user', 'url.count as 'cnt)
```

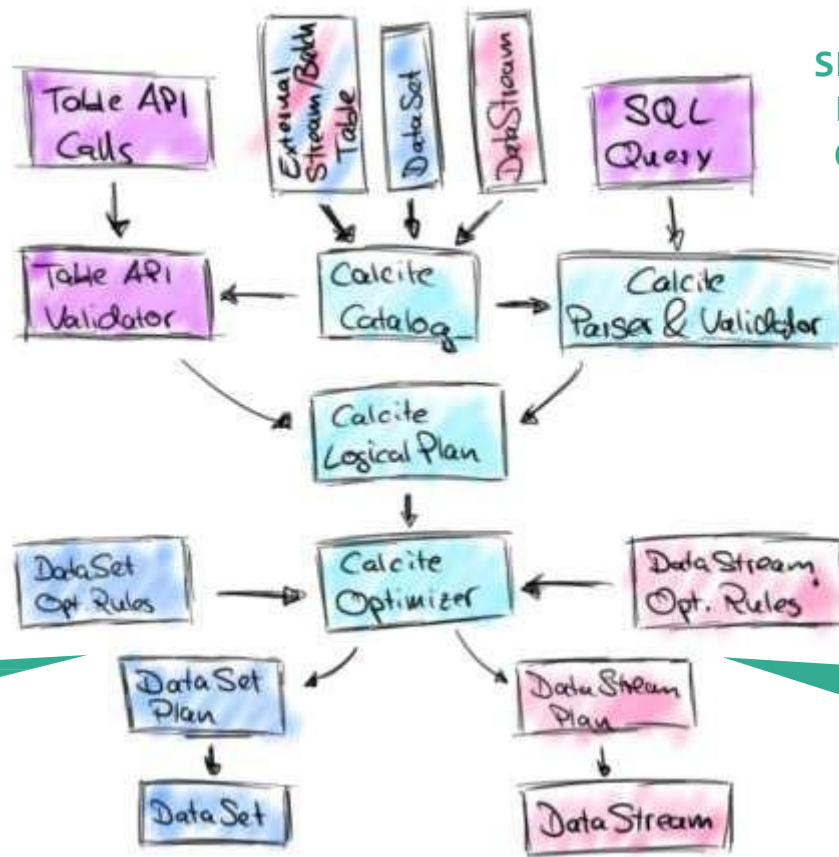
Unified APIs for batch & streaming data

*A query specifies exactly the same result  
regardless whether its input is  
static batch data or streaming data.*

# Query Translation

```
tableEnvironment  
  .scan("clicks")  
  .groupBy('user')  
  .select('user', 'url.count')
```

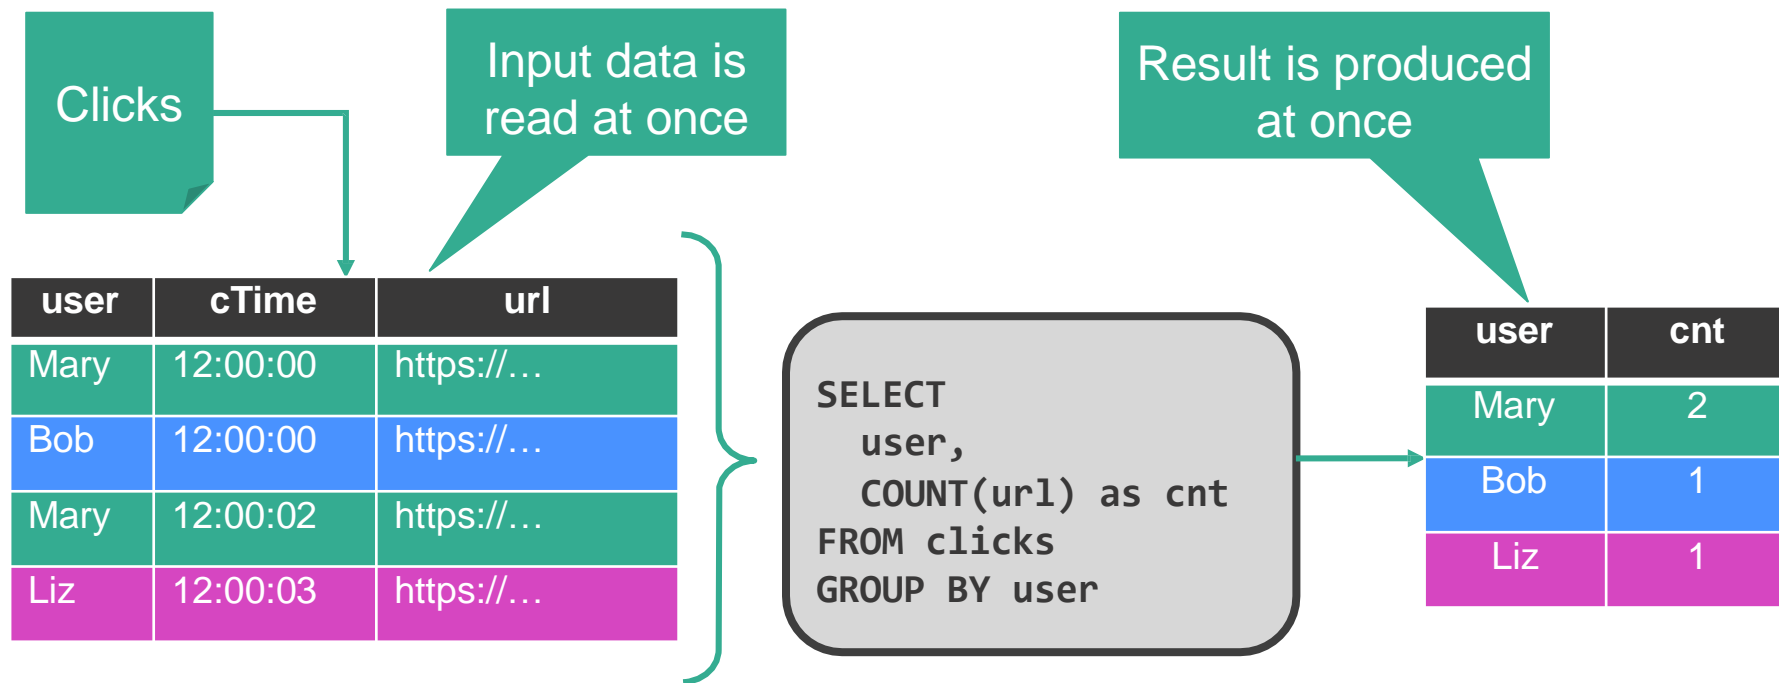
```
SELECT user, COUNT(url)  
FROM clicks  
GROUP BY user
```



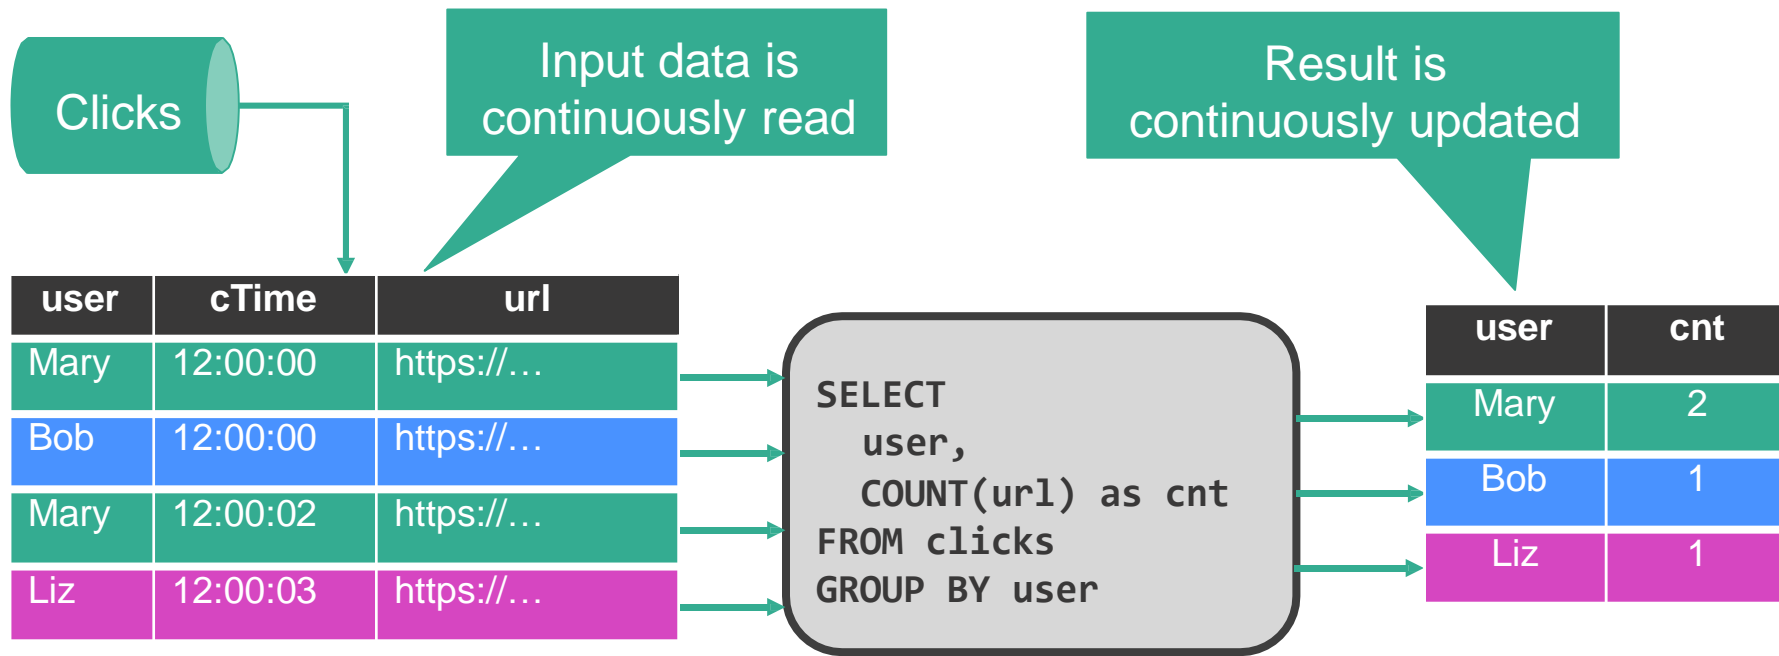
Input data is  
**bounded**  
(batch)

Input data is  
**unbounded**  
(streaming)

# What if “Clicks” is a File?



# What if “Clicks” is a Stream?



**The result is the same!**


# Why is Stream-Batch Unification Important?

- Usability

- ANSI SQL syntax: No custom “StreamSQL” syntax.
- ANSI SQL semantics: No stream-specific result semantics.

- Portability

- Run the same query on *bounded* & *unbounded* data
- Run the same query on *recorded* & *real-time* data
- *Bootstrapping* query state or *backfilling* results from historic data






# Database Systems Run Queries on Streams

- Materialized views (MV) are similar to regular views, but persisted to disk or memory
  - Used to speed-up analytical queries
  - MVs need to be updated when the base tables change
- MV maintenance is very similar to SQL on streams
  - Base table updates are a stream of DML statements
  - MV definition query is evaluated on that stream
  - MV is query result and continuously updated

# Continuous Queries in Flink

- Core concept is a “*Dynamic Table*”
  - Dynamic tables are changing over time
- Queries on dynamic tables
  - produce new dynamic tables (which are updated based on input)
  - do not terminate
- Stream  $\leftrightarrow$  Dynamic table conversions




# Stream ↔ Dynamic Table Conversions

- A stream is the changelog of a dynamic table
  - As change messages are ingested from a stream, a table evolves
  - As a table evolves, change messages are emitted to a stream
- Different changelog interpretations
  - Append-only change messages
  - Upsert change messages
  - Add/Retract change messages

# Continuous Queries

Based on the nature of the query, the result table might contain insert, update, upsert or delete records

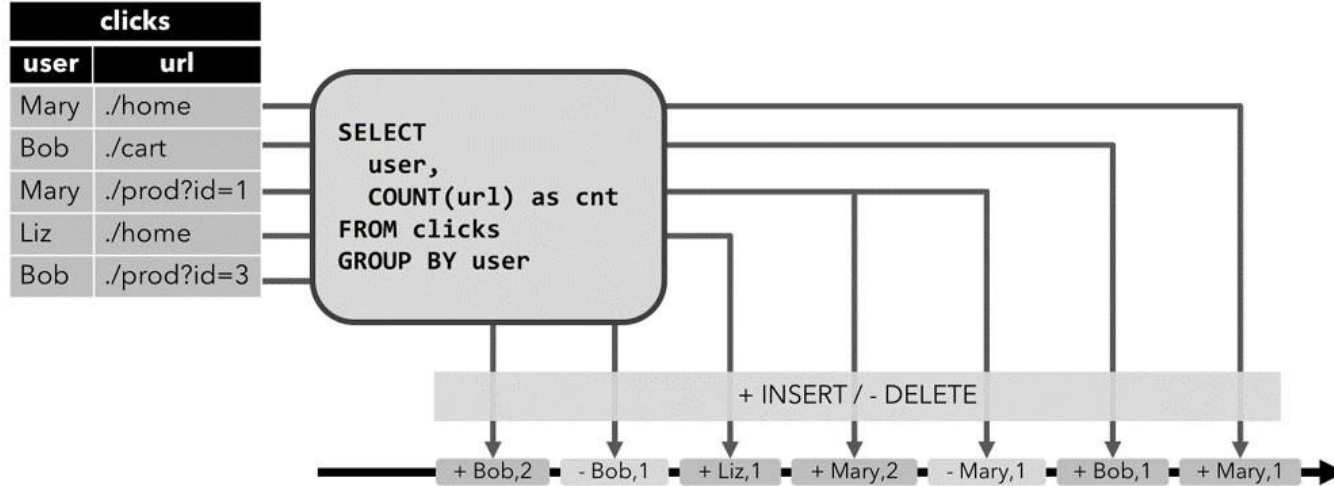


source: [https://ci.apache.org/projects/flink/flink-docs-release-1.7/dev/table/streaming/dynamic\\_tables.html](https://ci.apache.org/projects/flink/flink-docs-release-1.7/dev/table/streaming/dynamic_tables.html)

# Table to Stream (Concepts)

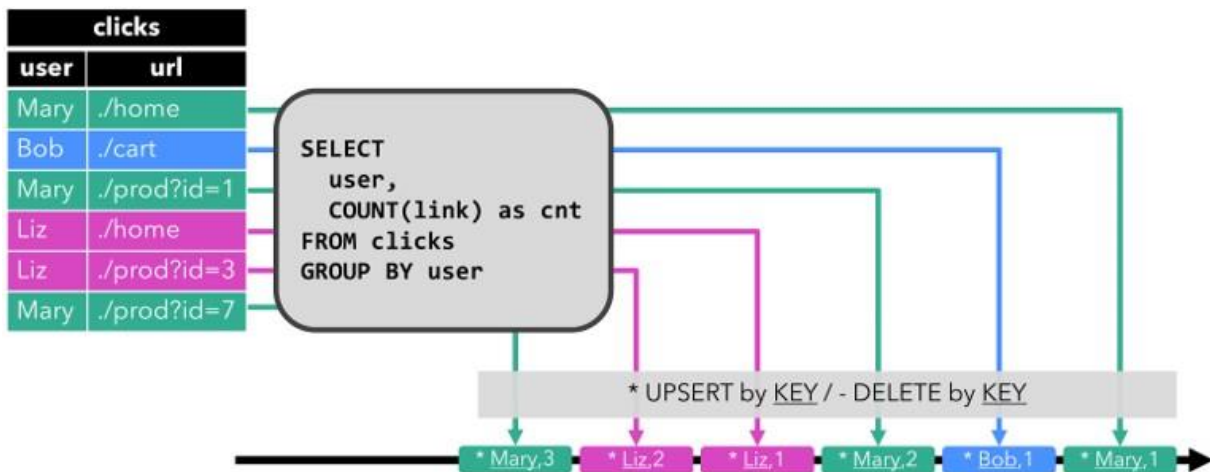
- When writing a dynamic table result back to a stream, we have to encode the type of change that occurred to the table: insert, delete or update
- Flink supports three ways to encode changes, recall the Dataflow model,
  - Append-only streams: only insert records are emitted on the stream
  - Retract streams: two messages are sent add and retract
    - Add: add message
    - Update: retract old value then add the new value
    - Delete: retract message
  - Upsert streams: both insert and update are handled as an upsert message. Delete is a retract message. Not yet implemented in Flink.

# Retract Stream



# Upsert Stream

- Less number of messages to be emitted
- Receiver operator has to be aware of message encoding
- A unique key is required



# Table API/SQL Examples

```
Table orders = tEnv.scan("Orders");  
// schema (a, b, c, rowtime)  
Table result = orders .filter("a.isNotNull && b.isNotNull && c.isNotNull")  
  .select("a.lowerCase() as a, b, rowtime")  
  .window(Tumble.over("1.hour").on("rowtime").as("w"))  
  .groupBy('w, a')  
  .select("a, w.end as hour, b.avg as avgBillingAmount");
```

```
SELECT A, AVG(B), TUMBLE_END(rowtime, INTERVAL '1' HOUR) as w.end  
FROM Orders GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), A  
WHERE A IS NOT NULL AND B IS NOT NULL AND C IS NOT NULL
```



# Table API Examples

```
Table orders = tEnv.scan("Orders");  
// schema (a, b, c, rowtime)  
Table result = orders .filter("a.isNotNull && b.isNotNull && c.isNotNull")  
    .select("a.lowerCase() as a, b, rowtime")  
    .window(Slide.over("10.minutes").every("5.minutes").on("rowtime").as("w"))  
    .groupBy('w, a')  
    .select("a, w.end as hour, b.avg as avgBillingAmount");
```

```
SELECT A, AVG(B), HOP_END(rowtime, INTERVAL '10' MINUTES, INTERVAL '5' MINUTES ) as w.end  
FROM Orders GROUP BY HOP(rowtime, INTERVAL '10' MINUTES, INTERVAL '5' MINUTES ), A  
WHERE A is NOT NULL AND B IS NOT NULL AND C IS NOT NULL
```

# Table API Examples

```
Table orders = tEnv.scan("Orders");  
// schema (a, b, c, rowtime)  
Table result = orders .filter("a.isNotNull && b.isNotNull && c.isNotNull")  
    .select("a.lowerCase() as _a, b, _rowtime")  
    .window(Slide.over("10.rows").every("5.rows").on("proctime").as("w"))  
    .groupBy("w, _a")  
    .select("a, _w.end as _hour, b.avg as _avgBillingAmount");
```

# Table API Examples

```
Table orders = tEnv.scan("Orders");  
// schema (a, b, c, rowtime)  
Table result = orders .filter("a.isNotNull && b.isNotNull && c.isNotNull")  
  .select("a.lowerCase() _as _a, _b, _rowtime")  
  .window(Session.withGap("10.minutes").on("rowtime").as("w"))  
  .groupBy("w, _a")  
  .select("a, _w.end _as _hour, _b.avg _as _avgBillingAmount");
```

```
SELECT A, AVG(B), SESSION_END(rowtime, INTERVAL '10' MINUTES) as w.end  
FROM Orders GROUP BY SESSION(rowtime, INTERVAL '10' MINUTES), A  
WHERE A IS NOT NULL AND B IS NOT NULL AND C IS NOT NULL
```

# SQL Feature Set in Flink $\geq$ 1.8.0

## STREAMING & BATCH

- SELECT FROM WHERE
- GROUP BY [HAVING]
  - Non-windowed
  - TUMBLE, HOP, SESSION windows
- JOIN
  - Time-Windowed INNER + OUTER JOIN
  - Non-windowed INNER + OUTER JOIN
- User-Defined Functions
  - Scalar
  - Aggregation
  - Table-valued

## STREAMING ONLY

- OVER / WINDOW
  - UNBOUNDED / BOUNDED PRECEDING
- INNER JOIN with time-versioned table
- MATCH\_RECOGNIZE
  - Pattern Matching/CEP (SQL:2016)


## BATCH ONLY

- UNION / INTERSECT / EXCEPT
- ORDER BY

# What Can I Build With That?


- Data Pipelines & Low-latency ETL

- Transform, aggregate, and move events in real-time
- Write streams to file systems, DBMS, K-V stores, ...
- Ingest appearing files to produce streams



- Stream & Batch Analytics

- Run analytical queries over bounded and unbounded data
- Query and compare historic and real-time data

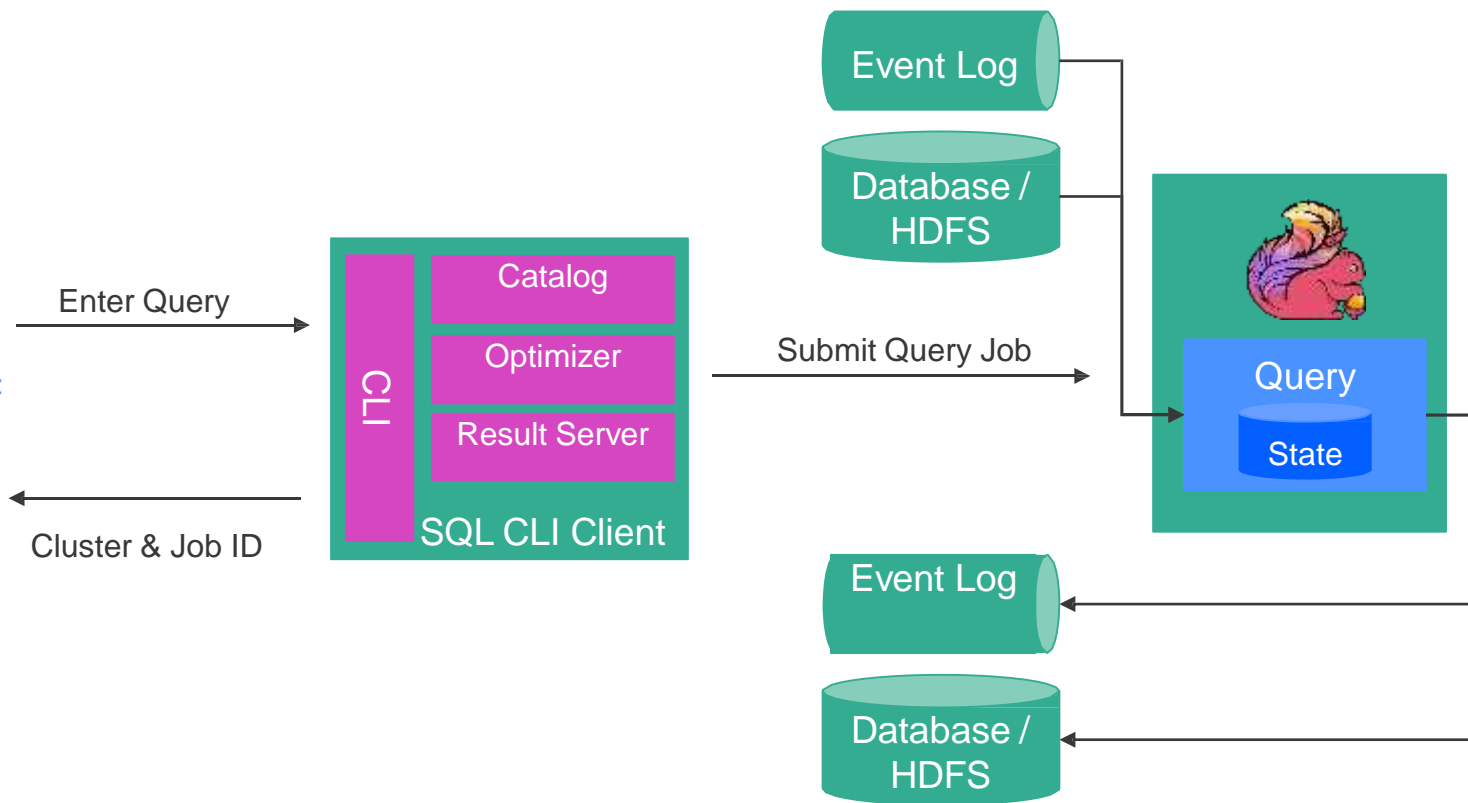


- Power Live Dashboards

- Compute and update data to visualize in real-time

# SQL CLI Client – Detached Queries

```
INSERT INTO  
sinkTable  
SELECT  
  user,  
  COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```



# The New York Taxi Rides Data Set

- A public data set about taxi rides in New York City
- Rides are ingested as append-only (streaming) table
  - Each ride is represented by a start and an end event

- **Table: Rides**

<code>rideId:</code>	<code>BIGINT</code>	<code>// ID of the taxi ride</code>
<code>taxiId:</code>	<code>BIGINT</code>	<code>// ID of the taxi</code>
<code>isStart:</code>	<code>BOOLEAN</code>	<code>// flag for pick-up (true) or drop-off (false) event</code>
<code>lon:</code>	<code>DOUBLE</code>	<code>// longitude of pick-up or drop-off location</code>
<code>lat:</code>	<code>DOUBLE</code>	<code>// latitude of pick-up or drop-off location</code>
<code>rideTime:</code>	<code>TIMESTAMP</code>	<code>// time of pick-up or drop-off event</code>
<code>psgCnt:</code>	<code>INT</code>	<code>// number of passengers</code>

# Compute Basic Statistics

- *Count rides* per *number of passengers*.

```
SELECT
    psgCnt,
    COUNT(*) as cnt
FROM Rides
WHERE isStart
GROUP BY
    psgCnt
```



# Identify Popular Pick-Up / Drop-Off Locations

- Compute *every 5 minutes* for *each area* the *number of departing and arriving taxis*.

```
SELECT
  area,
  isStart,
  TUMBLE_END(rideTime, INTERVAL '5' MINUTE) AS cntEnd,
  COUNT(*) AS cnt
FROM (SELECT rideTime, isStart, toAreaId(lon, lat) AS area
      FROM Rides)
GROUP BY
  area,
  isStart,
  TUMBLE(rideTime, INTERVAL '5' MINUTE)
```

# Average Tip Per Hour of Day

- Compute the *average tip* per *hour of day*. Fare data is stored in a separate table *Fares* that needs to be *joined*.

```
SELECT
    HOUR(r.rideTime) AS hourOfDay,
    AVG(f.tip) AS avgTip
FROM
    Rides r,
    Fares f
WHERE
    NOT r.isStart AND
    r.rideId = f.rideId AND
    f.payTime BETWEEN r.rideTime - INTERVAL '5' MINUTE AND r.rideTime
GROUP BY
    HOUR(r.rideTime);
```


# Average Ride Duration Per Pick-Up Location

- *Join start ride* and *end ride* events *on rideId* and compute *average ride duration per pick-up location*.

```
SELECT pickUpArea,  
       AVG(timeDiff(s.rowTime, e.rowTime) / 60000) AS avgDuration  
FROM (SELECT rideId, rowTime, toAreaId(lon, lat) AS pickUpArea  
      FROM TaxiRides  
      WHERE isStart) s  
JOIN  
      (SELECT rideId, rowTime  
      FROM TaxiRides  
      WHERE NOT isStart) e  
ON s.rideId = e.rideId AND  
   e.rowTime BETWEEN s.rowTime AND s.rowTime + INTERVAL '1' HOUR  
GROUP BY pickUpArea
```

# Number of rides per 30 minutes per area

Number of  
rides



# SQL Example

SELECT

toArealId(lat, lon) as cellId,

COUNT(distinct rideId) as rideCount,

TUMBLE\_ROWTIME(rowTime, INTERVAL '30' minute) AS rowTime,

TUMBLE\_START(rowTime, INTERVAL '30' minute) AS startTime,

TUMBLE\_END(rowTime, INTERVAL '30' minute) AS endTime

FROM

TaxiRides

GROUP BY

toArealId(lat, lon),

TUMBLE(rowTime, INTERVAL '30' minute)

# What is hard with SQL?

- Find rides with mid-stops



# Mid Stops: Pure SQL

## Rides table

<code>rideId:</code>	<code>BIGINT</code>	<code>// ID of the taxi ride</code>
<code>taxiId:</code>	<code>BIGINT</code>	<code>// ID of the taxi</code>
<code>isStart:</code>	<code>BOOLEAN</code>	<code>// flag for pick-up (true) or drop-off (false) event</code>
<code>lon:</code>	<code>DOUBLE</code>	<code>// longitude of pick-up or drop-off location</code>
<code>lat:</code>	<code>DOUBLE</code>	<code>// latitude of pick-up or drop-off location</code>
<code>rideTime:</code>	<code>TIMESTAMP</code>	<code>// time of pick-up or drop-off event</code>
<code>psgCnt:</code>	<code>INT</code>	<code>// number of passengers</code>

# Mid Stops: Pure SQL

```
select start.rideld from rides as start, rides as end
where start.taxild = end.taxild
and start.rideTime < end.rideTime
and start.isStart and end.isStart
and not exists (select 1 from rides as inbetweenNewRide
                where inbetweenNewRide.taxild = start.taxild
                and inbetweenNewRide.isStart
                and inbetweenNewRide.rideTime > start.rideTime
                and inbetweenNewRide.rideTime < end.rideTime      )
and not exists (select 1 from rides as inbetweenDrop
                where inbetweenDrop.taxild = start.taxild
                and inbetweenDrop.isStart = 0
                and inbetweenDrop.rideTime > start.rideTime
                and inbetweenDrop.rideTime < end.rideTime      )
```



# Mid Stops : Pattern Matching

```
Pattern.<Row>begin("S").where(  
  (row) -> {  
    return row.isStart == true;  
  }).next("E").where( (row) -> {  
    return row.isStart == true;  
  });
```

```
CEP.pattern(input.keyBy("driverId"),  
pattern)  
  .flatSelect(  
    new PatternFlatSelectFunction<Row,  
    Row>() {  
      @Override  
      public void flatSelect(  
        Map<String, List<Row>> pattern,  
        Collector<Row> out) throws Exception {  
        out.collect((  
          pattern.get("S").get(0).getRideId  
        ));  
      }  
    });
```

# MATCH\_RECOGNIZE

---

SQL:2016 extension

# Common use-cases

- stock market analysis
- customer behaviour
- tracking money laundering
- service quality
- network intrusion detection

# Position in a SQL Query

SELECT ...

FROM ...

**MATCH\_RECOGNIZE**

**(...)**

WHERE ...

GROUP BY ...

# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.ridId as sRidId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```



partition the data by  
given field = keyBy



# Rides with mid-stops


```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideld as sRideld  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```

specify order  
primary order = Event  
or Processing time



# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    S.rideId as sRideId  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (S E)  
  DEFINE  
    S AS S.isStart = true,  
    E AS E.isStart = true  
)
```



construct pattern

# PATTERN: Defining a Pattern

- *Concatenation:*
  - All rows of a pattern must be mapped to pattern variables
  - A pattern like (A B) means that the contiguity is strict between A and B
  - In other words: No rows between A and B
- *Quantifiers*
  - Number of rows mapped to a pattern variable


*	0 or more rows
+	1 or more rows
?	0 or 1 rows
{ n }, { n, }, { n, m }, { , m }	Define intervals (inclusive)
B*?	Perform mapping <i>reluctant</i> instead of <i>greedy</i> (default behavior)



# Rides with mid-stops

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
    PARTITION BY driverId  
    ORDER BY rowTime  
    MEASURES  
        S.rideld as sRideld  
  
    AFTER MATCH SKIP PAST LAST ROW  
    PATTERN (S E)  
    DEFINE  
        S AS S.isStart = true,  
        E AS E.isStart = true  
)
```

extract measures from  
matched sequence



# DEFINE/MEASURES: Define/Access Variables


- MEASURES

- Defines what will be included in the output of a matching pattern
- Project columns and define expressions for evaluation
- Number of produced rows depends on the output mode.  
Currently, `ONE ROW PER MATCH` = one output summary row per match only
- Output schema: *[partitioning columns] + [measures columns]*

- DEFINE

- Conditions that a row has to fulfill to be classified to the corresponding variable
- No condition for a variable evaluates to `TRUE`

# Multi-Stops



# Rides with more than one mid-stop

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
    PARTITION BY driverId  
    ORDER BY rowTime  
    MEASURES  
        S.rideId as sRideId  
    AFTER MATCH SKIP PAST LAST ROW  
    PATTERN (S E)  
    DEFINE  
        S AS S.isStart = true,  
        E AS E.isStart = true  
)
```

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
    PARTITION BY driverId  
    ORDER BY rowTime  
    MEASURES  
        S.rideId as sRideId,  
        COUNT(M.rideId) as countMidStops  
    AFTER MATCH SKIP PAST LAST ROW  
    PATTERN (S M{2,} E)  
    DEFINE  
        S AS S.isStart = true,  
        M AS M.rideId <> S.rideId,  
        E AS E.isStart = false AND E.rideId = S.rideId  
)
```


# Rides with more than one mid-stop

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
    PARTITION BY driverId  
    ORDER BY rowTime  
    MEASURES  
        S.rideld as sRideld  
    AFTER MATCH SKIP PAST LAST ROW  
    PATTERN (S E)  
    DEFINE  
        S AS S.isStart = true,  
        E AS E.isStart = true  
)
```

```
SELECT *  
FROM TaxiRides  
MATCH_RECOGNIZE (  
    PARTITION BY driverId  
    ORDER BY rowTime  
    MEASURES  
        S.rideld as sRideld,  
        COUNT(M.rideld) as countMidStops  
    AFTER MATCH SKIP PAST LAST ROW  
    PATTERN (S M{2,} E)  
    DEFINE  
        S AS S.isStart = true,  
        M AS M.rideld <> S.rideld,  
        E AS E.isStart = false AND E.rideld = S.rideld  
)
```

# Rush (peak) hours – V Shape Upside down

Number of rides



# Statistics per Area


```
CREATE VIEW RidesInArea AS
SELECT
    toAreald(lat, lon) as cellId,
    COUNT(distinct rideId) as rideCount,
    TUMBLE_ROWTIME(rowTime, INTERVAL '30' minute) AS rowTime,
    TUMBLE_START(rowTime, INTERVAL '30' minute) AS startTime,
    TUMBLE_END(rowTime, INTERVAL '30' minute) AS endTime
FROM
    TaxiRides
GROUP BY
    toAreald(lat, lon),
    TUMBLE(rowTime, INTERVAL '30' minute)
```

# Rush hours

```
SELECT * FROM RidesInArea  
MATCH_RECOGNIZE(  
  PARTITION BY cellId  
  ORDER BY rowTime  
  MEASURES  
    FIRST(UP.startTime) as rushStart,  
    LAST(DOWN.endTime) AS rushEnd,  
    SUM(UP.rideCount) + SUM(DOWN.rideCount) AS rideSum  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (UP{4,} DOWN{2,} E)  
  DEFINE  
    UP AS UP.rideCount > LAST(UP.rideCount, 1) or LAST(UP.rideCount, 1) IS NULL,  
    DOWN AS DOWN.rideCount < LAST(DOWN.rideCount, 1) OR  
      LAST(DOWN.rideCount, 1) IS NULL,  
    E AS E.rideCount > LAST(DOWN.rideCount)  
)
```

Use previous  
table/view

Number of rides





# Rush hours

SELECT \* FROM RidesInArea

MATCH\_RECOGNIZE(

PARTITION BY cellId

ORDER BY rowTime

MEASURES

FIRST(UP.startTime) as rushStart,

LAST(DOWN.endTime) AS rushEnd,

SUM(UP.rideCount) + SUM(DOWN.rideCount) AS rideSum

AFTER MATCH SKIP PAST LAST ROW

PATTERN (UP{4,} DOWN{2,} E)

DEFINE

UP AS UP.rideCount > LAST(UP.rideCount, 1) or LAST(UP.rideCount, 1) IS NULL,


DOWN AS DOWN.rideCount < LAST(DOWN.rideCount, 1) OR

LAST(DOWN.rideCount, 1) IS NULL,

E AS E.rideCount > LAST(DOWN.rideCount)

)

Number of rides



Use previous  
table/view

Apply match to the  
result of the inner  
query

# Rush hours

```
SELECT * FROM RidesInArea
```

```
MATCH_RECOGNIZE(
```

```
  PARTITION BY cellId
```

```
  ORDER BY rowTime
```

```
  MEASURES
```

```
    FIRST(UP.startTime) as rushStart,
```

```
    LAST(DOWN.endTime) AS rushEnd,
```

```
    SUM(UP.rideCount) + SUM(DOWN.rideCount) AS rideSum
```

```
AFTER MATCH SKIP PAST LAST ROW
```

```
PATTERN (UP{4,} DOWN{2,} E)
```

```
DEFINE
```

```
  UP AS UP.rideCount > LAST(UP.rideCount, 1) or LAST(UP.rideCount, 1) IS NULL,
```


```
  DOWN AS DOWN.rideCount < LAST(DOWN.rideCount, 1) OR
```

```
    LAST(DOWN.rideCount, 1) IS NULL,
```

```
  E AS E.rideCount > LAST(DOWN.rideCount)
```

```
)
```

Number of rides



access elements of  
looping pattern in  
measures

Access elements in  
looping pattern define

# Rush hours

```
SELECT * FROM RidesInArea
```

```
MATCH_RECOGNIZE(
```

```
  PARTITION BY cellId
```

```
  ORDER BY rowTime
```

```
  MEASURES
```

```
    FIRST(UP.startTime) as rushStart,
```

```
    LAST(DOWN.endTime) AS rushEnd,
```

```
    SUM(UP.rideCount) + SUM(DOWN.rideCount) AS rideSum
```

```
AFTER MATCH SKIP PAST LAST ROW
```

```
PATTERN (UP{4,} DOWN{2,} E)
```

```
DEFINE
```


```
  UP AS UP.rideCount > LAST(UP.rideCount, 1) or LAST(UP.rideCount, 1) IS NULL,
```

```
  DOWN AS DOWN.rideCount < LAST(DOWN.rideCount, 1) OR
```

```
    LAST(DOWN.rideCount, 1) IS NULL,
```

```
  E AS E.rideCount > LAST(DOWN.rideCount)
```

```
)
```



# DEFINE/MEASURES: Define/Access Variables

- *Pattern Variable Referencing*

- Access to the set of rows mapped to a particular pattern variable (so far)
- `A.price` = set of rows mapped so far to A plus the current row if we try to match the current row to A
- If A is a set, the last row is selected for scalar operations.
- If no pattern variable is specified (e.g. `SUM(price)`), the default pattern variable "\*" is used. This set contains all rows matched for pattern + current row.

PATTERN (A B+)

DEFINE

A **AS** A.price > 10,

B **AS** B.price > A.price **AND**

**SUM(price) < 100 AND SUM(B.price) < 80**

# DEFINE/MEASURES: Define/Access Variables

- *Pattern Variable Navigation*

- Logical offsets enable navigation within the events that were mapped to a particular pattern variable.
- `FIRST(variable.field, n)`      `n` starts from the beginning
- `LAST(variable.field, n)`      `n` starts from the end

PATTERN (A B+)

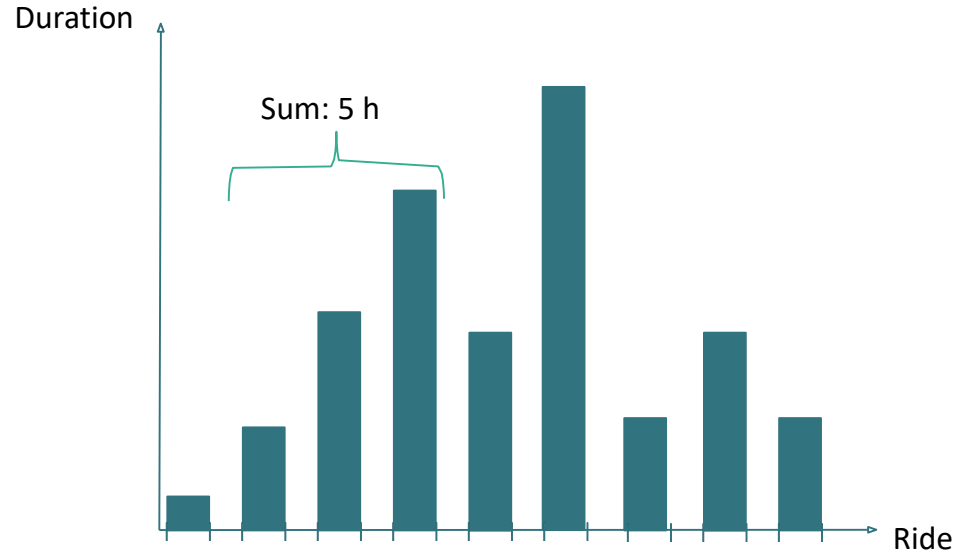
DEFINE

A **AS** A.price > 10,

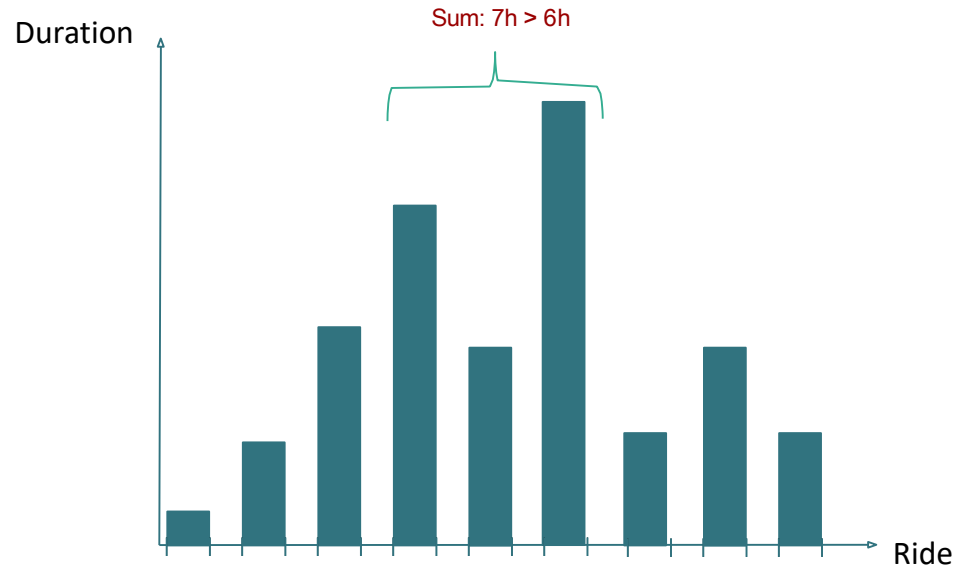
B **AS** (LAST(B.price, 1) **IS NULL OR** B.price > LAST(B.price, 1)) **AND**  
(LAST(B.price, 2) **IS NULL OR** B.price > 2 \* LAST(B.price, 2))

- Expressions on same "list" are supported: `LAST(A.price * A.tax)`

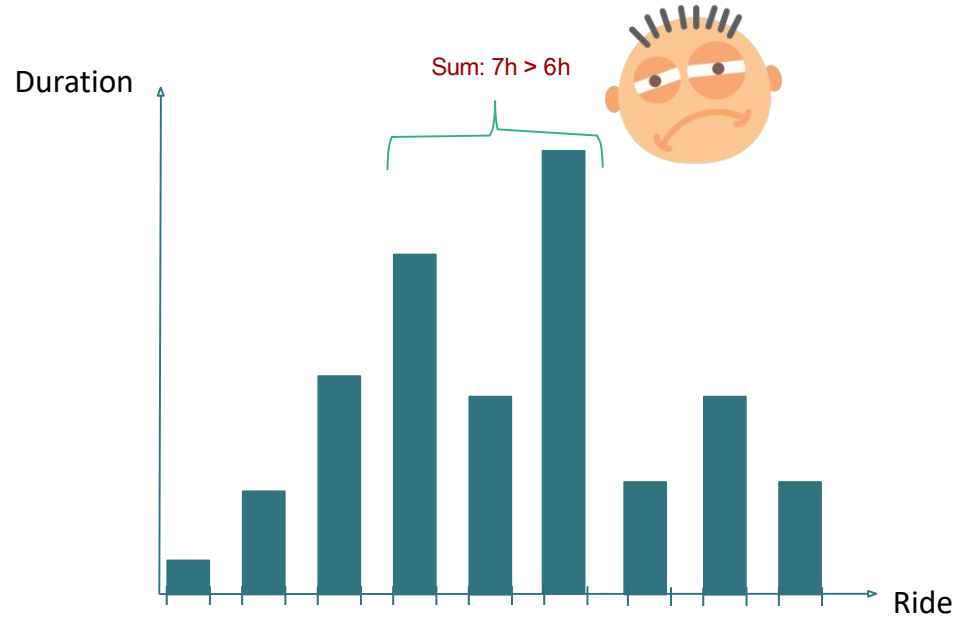
# Driver Fatigue



# Driver Fatigue



# Driver Fatigue



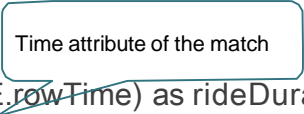


# Rides durations

```
CREATE VIEW RidesDurations AS
SELECT * FROM TaxiRides
MATCH_RECOGNIZE (
  PARTITION BY driverId
  ORDER BY rowTime
  MEASURES
    rideId as rideId,
    timeDiff(S.rowTime, E.rowTime) as rideDuration,
    MATCH_ROWTIME() as rowTime,
    S.rowTime as startTime,
    E.rowTime AS endTime
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (S E)
  DEFINE
    S AS S.isStart = true,
    E AS E.isStart = false AND E.rideId = S.rideId
);
```

# Rides durations

```
CREATE VIEW RidesDurations AS
SELECT * FROM TaxiRides
MATCH_RECOGNIZE (
  PARTITION BY driverId
  ORDER BY rowTime
  MEASURES
    rideId as rideId,
    timeDiff(S.rowTime, E.rowTime) as rideDuration,
    MATCH_ROWTIME() as rowTime,
    S.rowTime as startTime,
    E.rowTime AS endTime
  AFTER MATCH SKIP PAST LAST ROW
  PATTERN (S E)
  DEFINE
    S AS S.isStart = true,
    E AS E.isStart = false AND E.rideId = S.rideId
);
```



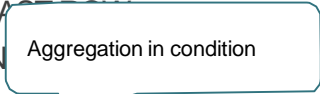
Time attribute of the match

# Drivers Fatigue

```
SELECT *  
FROM RidesDurations  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    formatDuration(SUM(rideDuration)) as totalRideDuration,  
    FIRST(R.startTime) as startTime,  
    LAST(R.endTime) as endTime  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (R+? E) WITHIN INTERVAL '1' DAY  
  DEFINE  
    E AS SUM(rideDuration) >= durationOfHours(2)  
);
```

# Drivers Fatigue

```
SELECT *  
FROM RidesDurations  
MATCH_RECOGNIZE (  
  PARTITION BY driverId  
  ORDER BY rowTime  
  MEASURES  
    formatDuration(SUM(rideDuration)) as totalRideDuration,  
    FIRST(R.startTime) as startTime,  
    LAST(R.endTime) as endTime  
  AFTER MATCH SKIP PAST LAST ROW  
  PATTERN (R+? E) WITHIN INTERVAL 1 HOUR  
  DEFINE  
    E AS SUM(rideDuration) >= durationOfHours(2)  
);
```



# Features set of MATCH\_RECOGNIZE

- Quantifiers support:
  - + (one or more), \* (zero or more), {x,y} (times)
  - greedy(default), ?(reluctant)
    - with some restrictions (not working for last pattern)
- After Match Skip
  - skip\_to\_first/last, skip\_past\_last, skip\_to\_next
- Aggregates (since 1.8)
- Allow time attribute extraction (since 1.8)
- Not supported:
  - alter(), permute, exclude '{- -}'

# AFTER MATCH SKIP: Continuation strategy

- Location where to start a new matching procedure after a complete match was found

SKIP PAST LAST ROW	next row after the last row of the current match
SKIP TO NEXT ROW	next row after the starting row of the match
SKIP TO LAST variable	last row that is mapped to the specified pattern variable
SKIP TO FIRST variable	first row that is mapped to the specified pattern variable

- Thus, also specifies how many matches a single event can belong to